

Settlement: The One Problem to Solve — and the Five We're Not

From: Product Management To: Greenroom Leadership Date: May 21, 2026

Team,

Settlement is a record problem before it's a math problem — and really six distinct problems before it's either.

Recommendation

I'd suggest we ship one thing this cycle: the **append-only audit ledger** — and deliberately hold the calculator and the walkthrough. All three are severe pain points, and all three are coming; the question is only what goes first. My case is that shipping one of them well will do more than shipping three half-built. Since the goal we've set is artist trust (sponsors follow artists), and settlement is where that trust is won or lost, the audit trail looks like the leverage point: it's the one piece that addresses trust failures directly rather than saving time around them. I'll lay out the full ranking and, just as importantly, defend what I'm choosing not to do.

Settlement is six problems wearing one name

"Fix settlement" sounds like one task. It's really six, and they differ in dependencies, effort, and payoff — which is how a team can end up with a polished calculator sitting on top of a record nobody can reconstruct. As Mariana, a booker, put it: "the actual settlement is a conversation, not a calculation." Here's how I'd rank the pieces — by leverage toward trust and by ROI — with the line drawn after the first.

1. **Audit trail — highest leverage on trust. Shipping now.** Trust breaks on opaque history, not arithmetic. Sarah Kim (WME) described a deal with three conflicting private versions — "the deal was a ghost"; Mariana's complaint was that "there's no record of what the tour manager actually said yes to." The Coastal Spell dispute (\$720, March 2025) came down to four ambiguous words in a deal email — both sides' math was correct. A ledger with evidence and ambiguity flags is the one item here aimed squarely at that failure. Moderate effort, and it lands where trust actually breaks: the deal stage, where the ambiguity originates. (It won't capture the final settlement math on its own — more on that cost below.)
2. **Calculator — highest ROI. Immediately next.** Deal types are already modeled; a booker can enter a vs-deal or a percentage-of-net. What's missing is computation — the app won't settle them, so deal modeling is only half-finished. Mariana tried the tool in 2023, hit deals it couldn't handle, and "now I just do it in the sheet." Roughly 56% of deals are structures the tool can't compute — so more than half of all settlements happen outside our product. The payoff here is real but mostly a booker one: it ends the daily spreadsheet tax and pulls that work back in-app. Its effect on artist trust is indirect — it surfaces only once the walkthrough exists to show the agent the math. Enormous ROI, immediately next, but it earns its trust through the walkthrough rather than on its own, which is why it follows the trail.
3. **2 a.m. walkthrough — high value, fully blocked. Third.** The traceable timeline that walks a tour manager from deal to payout is what Sarah ("if the settlement statement showed its work...") and Diego ("look at the math before I sit down") are asking for. It's genuinely worth building — and it can't do anything until 1 and 2 exist, since all it does is present their output. So it's sequenced behind them, not cut on merit.
4. **Real-time prediction — important, not leverage. Later.** Marcus, the GM, wants to forecast overage expenses and ticket sales for budgeting, and he's right that it matters — knowing the numbers before the show would genuinely help venues plan. But it's research-heavy, uncertain, and really several problems at once (expense forecasting isn't sales forecasting). It unblocks nothing above it and won't move trust this cycle, so I'd hold it: high importance, low leverage for now.
5. **Post-show agent comms & dispute resolution — downstream. Last.** Both matter, but they rest on the top three. Once those exist, much of this resolves on its own — and when a dispute does surface, the booker can handle it with what's shipping now: export the ledger and show the agent the record. Once the timeline lands, that becomes a walkthrough too. The record is already most of the dispute tool, so building dedicated comms or arbitration features ahead of its foundation would be low marginal return.

What “ship the audit trail” actually means

Append-only ledger. Every show gets one chronological record. Each entry captures timestamp, author, free-text note, optional evidence, and any structured field change it drove. Entries are immutable — corrections are new entries, so a \$500→\$700 hospitality cap shows both states and why it moved.

AI summary. A readable state view on top of the ledger, every line expandable to its source entry. Strictly assistive — when summary and ledger disagree, the ledger governs, and the product says so.

Smart Companion. Proposes, the human confirms, and only then it acts. Flags field/note contradictions, conflicting states (disputed show with a positive sign-off), and dispute-prone ambiguity (“we’ll probably cover hospitality”) with a neutral prompt — clarity, not leverage. Evidence stays light: a screenshot, a pasted message, or “confirmed verbally on the call.”

The cut I’m defending

Holding the calculator costs us something real, and I don’t want to gloss over it: the ledger’s **settlement-stage** entries stay thin until the calculator exists, because the final numbers still live in Mariana’s spreadsheet. I’m willing to accept that, and here’s the reasoning — the trail’s value concentrates in the **deal-evolution stage** (show creation, terms, the agent’s email, the ambiguity), which needs no calculator and is exactly where the trust failures start. So we capture the high-trust half now and feed the settlement half next cycle. Trying to ship both at once is what risks two half-built systems and dilutes the one bet that actually moves the goal. I’d rather do less, and do it well.

Risks

Over-formalization (the one I’d watch most). If logging ever feels like building a case against the other party, people will route around it and go off-platform — the very failure we’re trying to fix. Mitigation: evidence optional, note-only entries allowed, defaults kept light. **AI reliability.** A summary that misreads deal language creates false confidence — so the raw ledger stays prominent, every line is traceable, and suggestions sit behind a confidence threshold. **Fatigue.** More logging at 2 a.m. is a real cost; minimize required fields, keep entry to free text, avoid multi-step forms.

How we’ll know it worked

Leading signals are adoption: share of shows using the ledger, share with evidence attached, Companion-flag acceptance, and — the canary I’d watch hardest — whether bookers log in-app instead of exporting to a spreadsheet. Outcomes that matter: fewer disputed settlements, fewer post-show revisions, faster resolution, rising NPS from touring teams. Marcus spends ~25 hours a month on settlement and cleanup; the trail won’t reclaim all of it, but disputes tied to ambiguous terms should fall first.

Bottom line

One bet this cycle: the audit trail. It’s the piece that moves trust most directly, and it’s the foundation the calculator and the walkthrough both stand on. If we build the high-trust half well now, we earn the right — and a much better starting point — to build the rest. Happy to talk through any of these calls; the cuts especially are worth pressure-testing.

— Product Management